


Operating Systems

*Narasimhulu M*_{M. Tech.}
Assistant Professor
Department of Computer Science & Engineering



S.No.	Course Outcomes	Cognitive Level
1	Illustrate various types of system calls and find the stages of various process states.	Understand
2	Implement thread scheduling and process scheduling techniques	Apply
3	Distinguish among IPC synchronization Techniques	Understand
4	Implement page replacement algorithms, memory management techniques and deadlock issues.	Apply
5	Make use of the file systems for applying different allocation and access techniques.	Understand
6	Illustrate system protection and Security.	Understand



Unit-2: Process Management

- **Processes:** Process Concept, Scheduling, Operations. Inter process Communication: Shared-Memory Systems, Message-Passing Systems, Examples, Communication in Client-Server Systems. CPU Scheduling: Scheduling Criteria, Scheduling Algorithms, Threads.
- **Process Synchronization:** The critical-section problem, Petersons Solution, Synchronization Hardware, Mutex Locks, Semaphores, Classic problems of synchronization, Monitors.

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

3



Unit 2 – Process Management

Narasimhulu M. *M. Tech.*

Assistant Professor

Department of Computer Science & Engineering



Chapter 1

Process

Narasimhulu M_{M. Tech.}
Assistant Professor
Department of Computer Science & Engineering



Process Concept

Narasimhulu M_{M. Tech.}
Assistant Professor
Department of Computer Science & Engineering



Process Concept

- An operating system executes a variety of programs that run as a process.
- **Process** – a program in execution; process execution must progress in sequential fashion. No parallel execution of instructions of a single process
- Multiple parts
 - The program code, also called **text section**
 - Current activity including **program counter**, processor registers
 - **Stack** containing temporary data
 - Function parameters, return addresses, local variables
 - **Data section** containing global variables
 - **Heap** containing memory dynamically allocated during run time

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

7



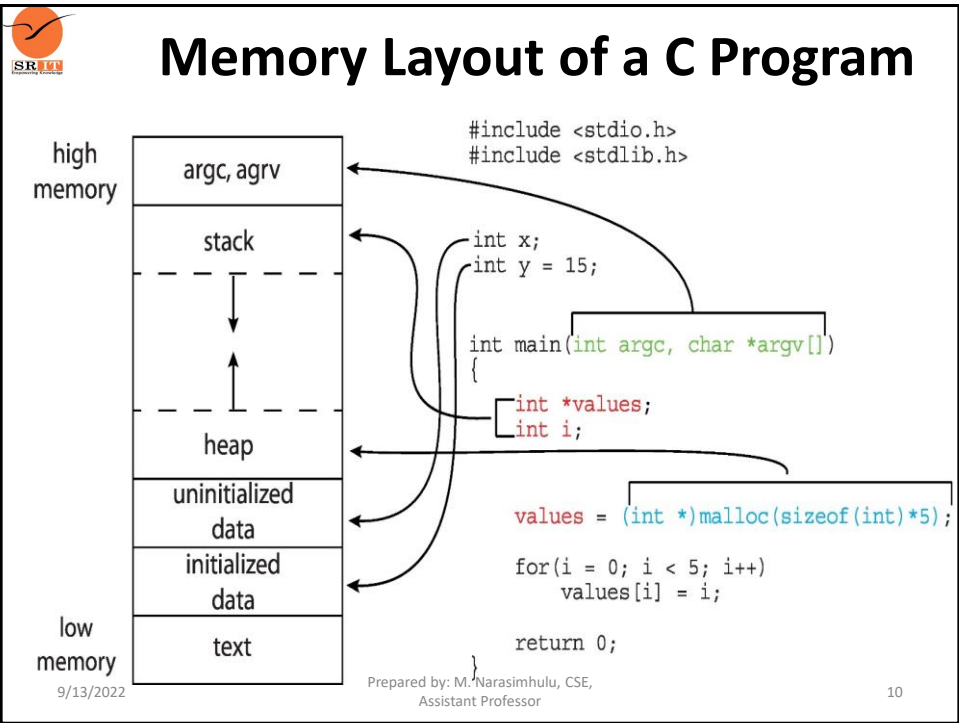
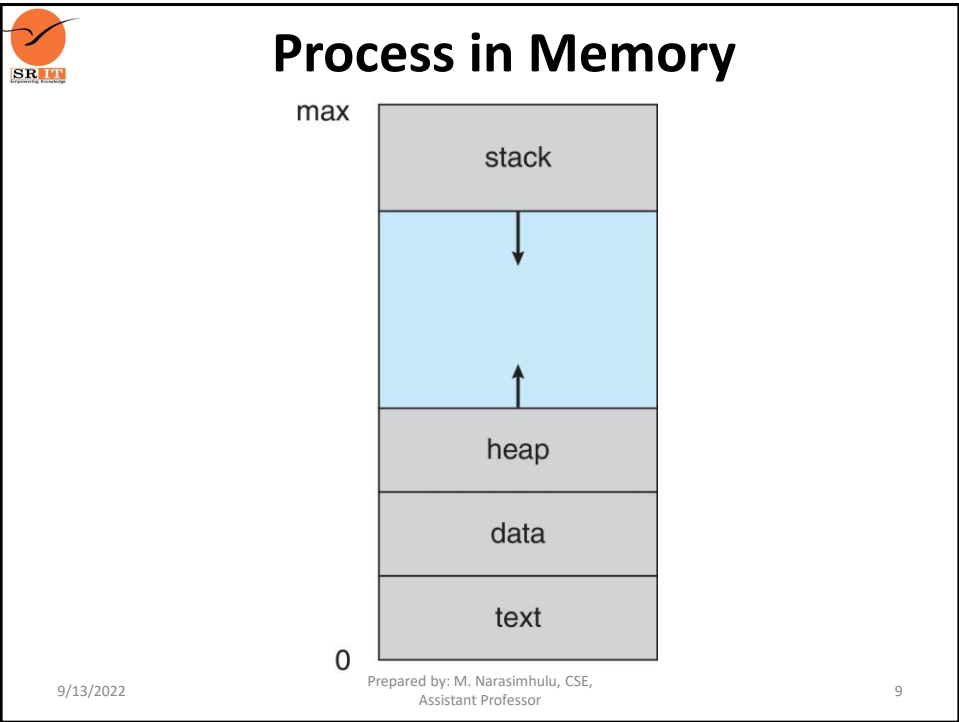
Process Concept (Cont.)

- Program is **passive** entity stored on disk (**executable file**); process is **active**
 - Program becomes process when an executable file is loaded into memory
- Execution of program started via GUI mouse clicks, command line entry of its name, etc.
- One program can be several processes
 - Consider multiple users executing the same program
 - Compiler
 - Text editor

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

8





Process State

- As a process executes, it changes **state**
 - **New**: The process is being created
 - **Running**: Instructions are being executed
 - **Waiting**: The process is waiting for some event to occur
 - **Ready**: The process is waiting to be assigned to a processor
 - **Terminated**: The process has finished execution

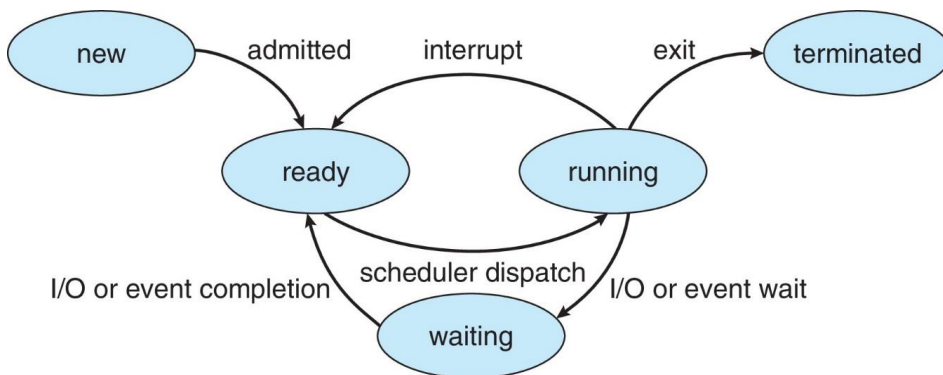
9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

11




Diagram of Process State



9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

12



Process Control Block (PCB)

Information associated with each process(also called **task control block**)


- Process state – running, waiting, etc.
- Program counter – location of instruction to next execute
- CPU registers – contents of all process-centric registers
- CPU scheduling information- priorities, scheduling queue pointers
- Memory-management information – memory allocated to the process
- Accounting information – CPU used, clock time elapsed since start, time limits
- I/O status information – I/O devices allocated to process, list of open files

process state
process number
program counter
registers
memory limits
list of open files
...

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

13




Threads

- So far, process has a single thread of execution
- Consider having multiple program counters per process
 - Multiple locations can execute at once
 - Multiple threads of control -> **threads**
- Must then have storage for thread details, multiple program counters in PCB
- Explore in detail in Chapter 4

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

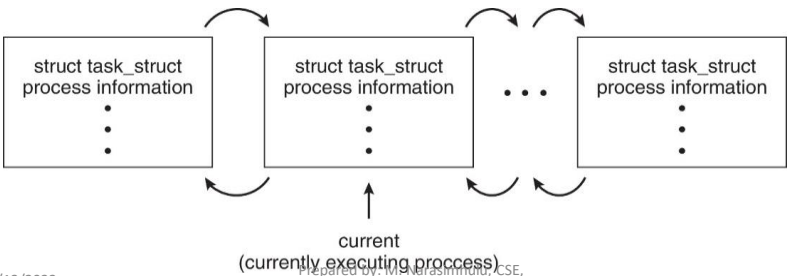
14



Process Representation in Linux

Represented by the C structure task_struct


```
pid_t pid; /* process identifier */
long state; /* state of the process */
unsigned int time_slice; /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm; /* address space of this process */
```



9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

15



Process Scheduling

Narasimhulu M.
M. Tech.
Assistant Professor
Department of Computer Science & Engineering



Process Scheduling

- **Process scheduler** selects among available processes for next execution on CPU core
- Goal -- Maximize CPU use, quickly switch processes onto CPU core
- Maintains **scheduling queues** of processes
 - **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
 - **Wait queues** – set of processes waiting for an event (i.e., I/O)
 - Processes migrate among the various queues

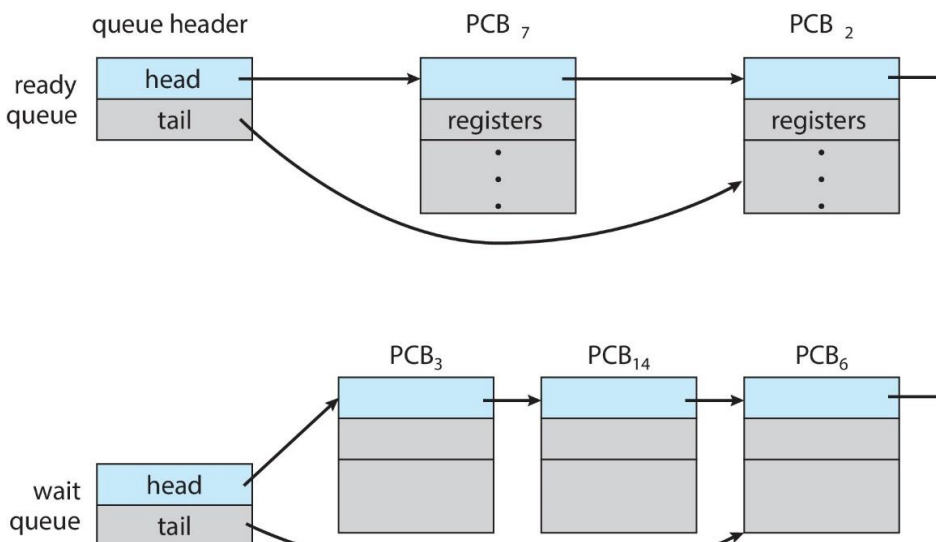
9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

17



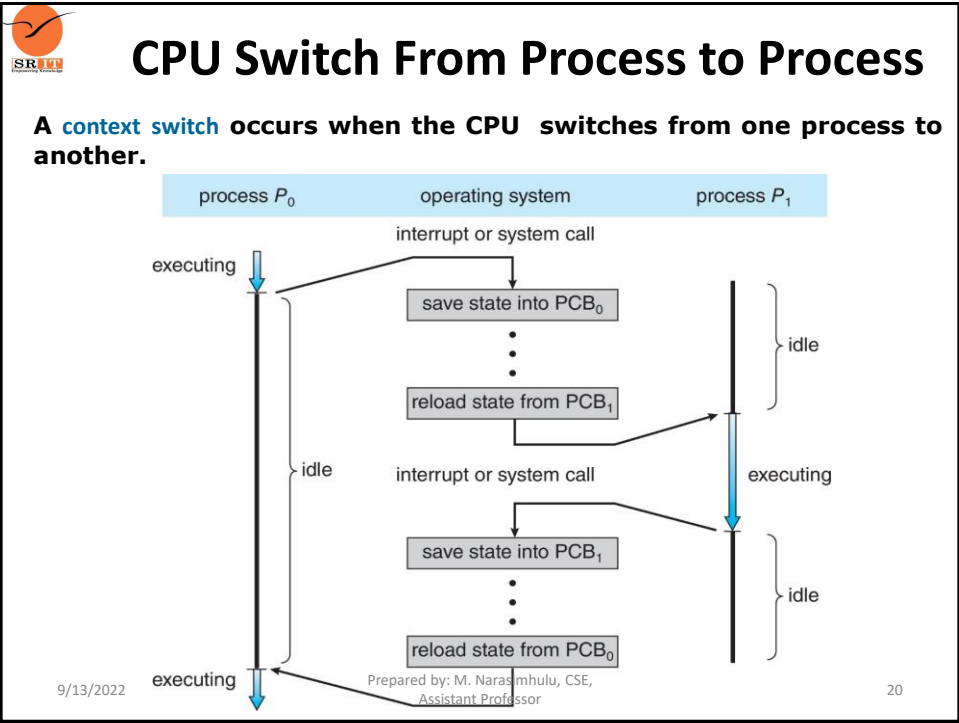
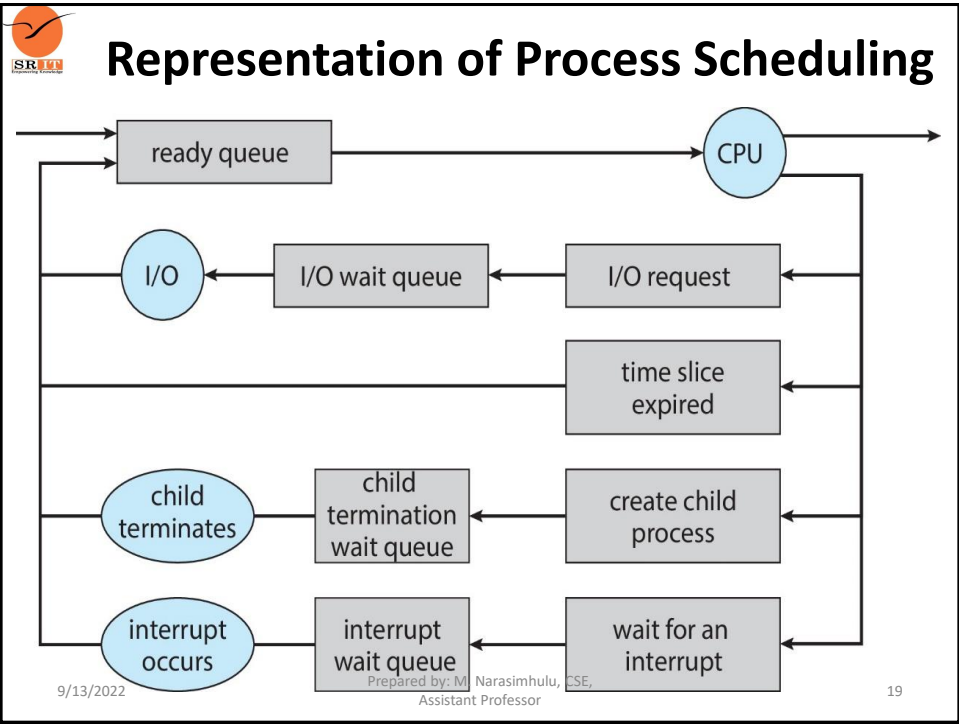
Ready and Wait Queues



9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

18





Context Switch

- When CPU switches to another process, the system must **save the state** of the old process and load the **saved state** for the new process via a **context switch**
- **Context** of a process represented in the PCB
- Context-switch time is pure overhead; the system does no useful work while switching
 - The more complex the OS and the PCB → the longer the context switch
- Time dependent on hardware support
 - Some hardware provides multiple sets of registers per CPU → multiple contexts loaded at once

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

21



Multitasking in Mobile Systems

- Some mobile systems (e.g., early version of iOS) allow only one process to run, others suspended
- Due to screen real estate, user interface limits iOS provides for a
 - Single **foreground** process- controlled via user interface
 - Multiple **background** processes– in memory, running, but not on the display, and with limits
 - Limits include single, short task, receiving notification of events, specific long-running tasks like audio playback
- Android runs foreground and background, with fewer limits
 - Background process uses a **service** to perform tasks
 - Service can keep running even if background process is suspended
 - Service has no user interface, small memory use

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

22



Operations on Processes

Narasimhulu M_{M. Tech.}

Assistant Professor

Department of Computer Science & Engineering



Operations on Processes

- System must provide mechanisms for:
 - Process creation
 - Process termination



Process Creation

- **Parent** process create **children** processes, which, in turn create other processes, forming a **tree** of processes
- Generally, process identified and managed via a **process identifier (pid)**
- Resource sharing options
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources
- Execution options
 - Parent and children execute concurrently
 - Parent waits until children terminate

9/13/2022


Prepared by: M. Narasimhulu, CSE,
Assistant Professor

25



Process Creation (Cont.)

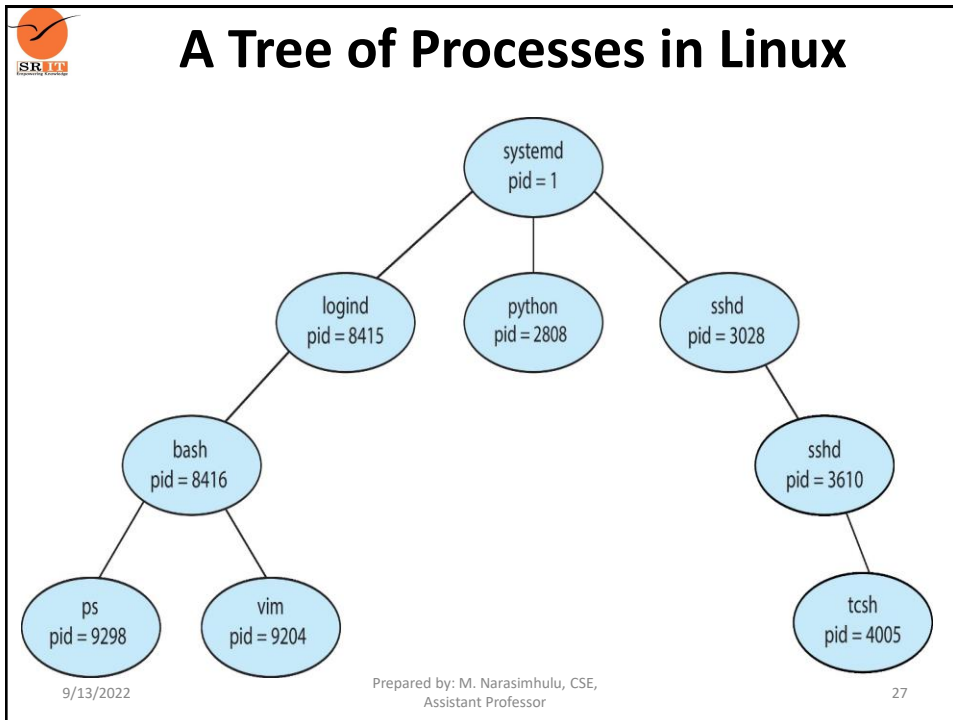
- Address space
 - Child duplicate of parent
 - Child has a program loaded into it
- UNIX examples
 - `fork()` system call creates new process
 - `exec()` system call used after a `fork()` to replace the process' memory space with a new program
 - Parent process calls `wait()` waiting for the child to terminate



9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

26



```

#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
  
```

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

28



Creating a Separate Process via Windows API

```
#include <stdio.h>
#include <windows.h>

int main(VOID)
{
    STARTUPINFO si;
    PROCESS_INFORMATION pi;

    /* allocate memory */
    ZeroMemory(&si, sizeof(si));
    si.cb = sizeof(si);
    ZeroMemory(&pi, sizeof(pi));

    /* create child process */
    if (!CreateProcess(NULL, /* use command line */
        "C:\\WINDOWS\\system32\\mspaint.exe", /* command */
        NULL, /* don't inherit process handle */
        NULL, /* don't inherit thread handle */
        FALSE, /* disable handle inheritance */
        0, /* no creation flags */
        NULL, /* use parent's environment block */
        NULL, /* use parent's existing directory */
        &si,
        &pi))
    {
        fprintf(stderr, "Create Process Failed");
        return -1;
    }
    /* parent will wait for the child to complete */
    WaitForSingleObject(pi.hProcess, INFINITE);
    printf("Child Complete");

    /* close handles */
    CloseHandle(pi.hProcess);
    CloseHandle(pi.hThread);
}
```

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

29



Process Termination

- Process executes last statement and then asks the operating system to delete it using the `exit()` system call.
 - Returns status data from child to parent (via `wait()`)
 - Process' resources are deallocated by operating system
- Parent may terminate the execution of children processes using the `abort()` system call. Some reasons for doing so:
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - The parent is exiting, and the operating systems does not allow a child to continue if its parent terminates

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

30



Process Termination

- Some operating systems do not allow child to exist if its parent has terminated. If a process terminates, then all its children must also be terminated.
 - **cascading termination.** All children, grandchildren, etc., are terminated.
 - The termination is initiated by the operating system.
- The parent process may wait for termination of a child process by using the `wait()` system call. The call returns status information and the pid of the terminated process

```
pid = wait(&status);
```

- If no parent waiting (did not invoke `wait()`) process is a **zombie**
- If parent terminated without invoking `wait()`, process is an **orphan**

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

31




Android Process Importance Hierarchy

- Mobile operating systems often have to terminate processes to reclaim system resources such as memory. From **most** to **least** important:
 - Foreground process
 - Visible process
 - Service process
 - Background process
 - Empty process
- Android will begin terminating processes that are least important.

9/13/2022

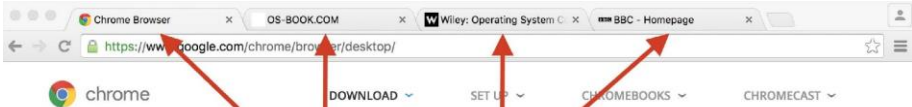
Prepared by: M. Narasimhulu, CSE,
Assistant Professor

32




Multiprocess Architecture – Chrome Browser

- Many web browsers ran as single process (some still do)
 - If one web site causes trouble, entire browser can hang or crash
- Google Chrome Browser is multiprocess with 3 different types of processes:
 - **Browser** process manages user interface, disk and network I/O
 - **Renderer** process renders web pages, deals with HTML, Javascript. A new renderer created for each website opened
 - Runs in **sandbox** restricting disk and network I/O, minimizing effect of security exploits
 - **Plug-in** process for each type of plug-in



Each tab represents a separate process.

9/13/2022
Prepared by: M. Narasimhulu, CSE,
Assistant Professor
33



Interprocess Communication

Narasimhulu M_{M. Tech.}

Assistant Professor

Department of Computer Science & Engineering



Interprocess Communication

- Processes within a system may be **independent** or **cooperating**
- Cooperating process can affect or be affected by other processes, including sharing data
- Reasons for cooperating processes:
 - Information sharing
 - Computation speedup
 - Modularity
 - Convenience
- Cooperating processes need **interprocess communication (IPC)**
- Two models of IPC
 - **Shared memory** (under the control of users)
 - **Message passing** (under the control of OS)

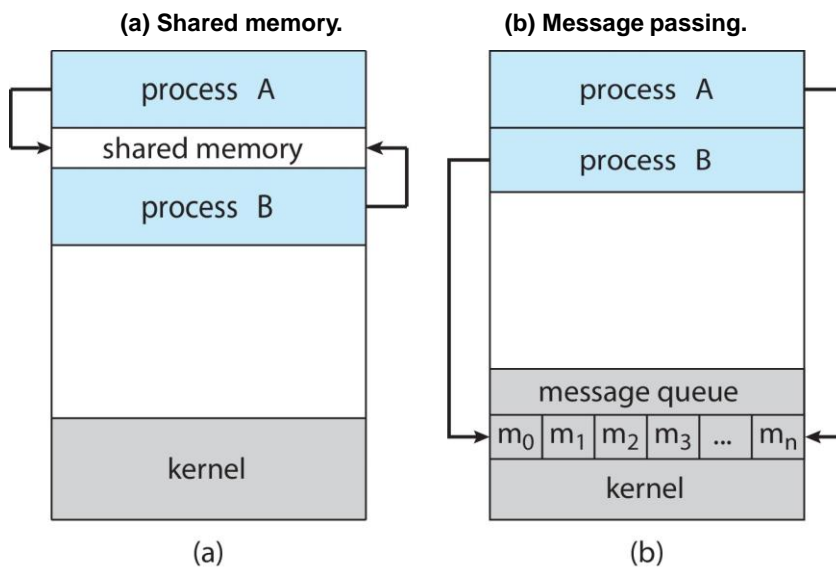
9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

35



Communications Models



9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

36



Producer-Consumer Problem

- Paradigm for cooperating processes:
 - **producer** process produces information that is consumed by a **consumer** process
- Two variations:
 - **unbounded-buffer** places no practical limit on the size of the buffer:
 - Producer never waits
 - Consumer waits if there is no buffer to consume
 - **bounded-buffer** assumes that there is a fixed buffer size
 - Producer must wait if all buffers are full
 - Consumer waits if there is no buffer to consume

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

37



Shared Memory Solution

- An area of memory shared among the processes that wish to communicate
- The communication is under the control of the users processes not the operating system.
- Major issues is to provide mechanism that will allow the user processes to synchronize their actions when they access shared memory.
- Synchronization is discussed in great details in Chapters 6 & 7.

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

38



Bounded-Buffer – Shared-Memory Solution

- Shared data

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```


- Solution presented in next slides is correct, but can only use $BUFFER_SIZE - 1$ items; that is: 9 items

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

39



Bounded-Buffer (Cont.)



9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

40



Producer Process – Shared Memory

```

item next_produced;

while (true) {
    /* produce an item in next produced */
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}

```



9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

41



Consumer Process – Shared Memory

```

item next_consumed;

while (true) {
    while (in == out)
        ; /* do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;

    /* consume the item in next
consumed */
}

```

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

42



What about Filling all the Buffers?

- Suppose that we wanted to provide a solution to the consumer-producer problem that fills **all** the buffers.
- We can do so by having an integer `counter` that keeps track of the number of full buffers.
- Initially, `counter` is set to 0.
- The integer `counter` is incremented by the producer after it produces a new buffer.
- The integer `counter` is and is decremented by the consumer after it consumes a buffer.

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

43



Producer

```
while (true) {
    /* produce an item in next produced */

    while (counter == BUFFER_SIZE)
        ; /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
    counter++;
}
```

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

44



Consumer

```
while (true) {
    while (counter == 0)
        ; /* do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    counter--;
    /* consume the item in next consumed */
}
```

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

45



Race Condition

- **counter++** could be implemented as

```
register1 = counter
register1 = register1 + 1
counter = register1
```

- **counter--** could be implemented as

```
register2 = counter
register2 = register2 - 1
counter = register2
```

- Consider this execution interleaving with “count = 5” initially:

S0: producer execute	register1 = counter	{register1 = 5}
S1: producer execute	register1 = register1 + 1	{register1 = 6}
S2: consumer execute	register2 = counter	{register2 = 5}
S3: consumer execute	register2 = register2 - 1	{register2 = 4}
S4: producer execute	counter = register1	{counter = 6}
S5: consumer execute	counter = register2	{counter = 4}

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

46



Race Condition (Cont.)

- Question – why was there no race condition in the first solution (where at most $N - 1$) buffers can be filled?
- More in Chapter 6.

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

47



IPC – Message Passing

- Processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
 - `send(message)`
 - `receive(message)`
- The *message* size is either fixed or variable

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

48



Message Passing (Cont.)

- If processes P and Q wish to communicate, they need to:
 - Establish a **communication link** between them
 - Exchange messages via send/receive
- Implementation issues:
 - How are links established?
 - Can a link be associated with more than two processes?
 - How many links can there be between every pair of communicating processes?
 - What is the capacity of a link?
 - Is the size of a message that the link can accommodate fixed or variable?
 - Is a link unidirectional or bi-directional?

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

49



Implementation of Communication Link

- Physical:
 - Shared memory
 - Hardware bus
 - Network
- Logical:
 - Direct or indirect
 - Synchronous or asynchronous
 - Automatic or explicit buffering

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

50



Direct Communication

- Processes must name each other explicitly:
 - `send(P, message)` – send a message to process P
 - `receive(Q, message)` – receive a message from process Q
- Properties of communication link
 - Links are established automatically
 - A link is associated with exactly one pair of communicating processes
 - Between each pair there exists exactly one link
 - The link may be unidirectional, but is usually bi-directional

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

51



Indirect Communication

- Messages are directed and received from mailboxes (also referred to as ports)
 - Each mailbox has a unique id
 - Processes can communicate only if they share a mailbox
- Properties of communication link
 - Link established only if processes share a common mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links
 - Link may be unidirectional or bi-directional

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

52



Indirect Communication (Cont.)

- Operations
 - Create a new mailbox (port)
 - Send and receive messages through mailbox
 - Delete a mailbox
- Primitives are defined as:
 - `send(A, message)` – send a message to mailbox A
 - `receive(A, message)` – receive a message from mailbox A

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

53



Indirect Communication (Cont.)

- Mailbox sharing
 - P_1 , P_2 , and P_3 share mailbox A
 - P_1 sends; P_2 and P_3 receive
 - Who gets the message?
- Solutions
 - Allow a link to be associated with at most two processes
 - Allow only one process at a time to execute a receive operation
 - Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

54



Synchronization

Message passing may be either blocking or non-blocking

- **Blocking** is considered **synchronous**
 - **Blocking send** -- the sender is blocked until the message is received
 - **Blocking receive** -- the receiver is blocked until a message is available
- **Non-blocking** is considered **asynchronous**
 - **Non-blocking send** -- the sender sends the message and continue
 - **Non-blocking receive** -- the receiver receives:
 - A valid message, or
 - Null message
- Different combinations possible
 - If both send and receive are blocking, we have a **rendezvous**

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

55



Producer-Consumer: Message Passing

- Producer


```
message next_produced;
while (true) {
    /* produce an item in next_produced */

    send(next_produced);
}
```
- Consumer


```
message next_consumed;
while (true) {
    receive(next_consumed)

    /* consume the item in next_consumed */
}
```

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

56



Buffering

- Queue of messages attached to the link.
- Implemented in one of three ways
 1. Zero capacity – no messages are queued on a link.
Sender must wait for receiver (rendezvous)
 2. Bounded capacity – finite length of n messages
Sender must wait if link full
 3. Unbounded capacity – infinite length Sender never waits

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

57



Examples

Narasimhulu M. *M. Tech.*

Assistant Professor

Department of Computer Science & Engineering



Examples of IPC Systems - POSIX

- POSIX Shared Memory

- Process first creates shared memory segment
`shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);`

- Also used to open an existing segment

- Set the size of the object

`ftruncate(shm_fd, 4096);`

- Use `mmap()` to memory-map a file pointer to the shared memory object

- Reading and writing to shared memory is done by using the pointer returned by `mmap()`.

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

59



IPC POSIX Producer

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "QS";
    /* strings written to shared memory */
    const char *message.0 = "Hello";
    const char *message.1 = "World!";

    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* create the shared memory object */
    shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);

    /* configure the size of the shared memory object */
    ftruncate(shm_fd, SIZE);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_WRITE, MAP_SHARED, shm_fd, 0);

    /* write to the shared memory object */
    sprintf(ptr, "%s", message.0);
    ptr += strlen(message.0);
    sprintf(ptr, "%s", message.1);
    ptr += strlen(message.1);

    return 0;
}
```

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

60



IPC POSIX Consumer

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* open the shared memory object */
    shm_fd = shm_open(name, O_RDONLY, 0666);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0);

    /* read from the shared memory object */
    printf("%s", (char *)ptr);

    /* remove the shared memory object */
    shm_unlink(name);

    return 0;
}
```

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

61



Examples of IPC Systems - Mach

- Mach communication is message based
 - Even system calls are messages
 - Each task gets two ports at creation- Kernel and Notify
 - Messages are sent and received using the **mach_msg()** function
 - Ports needed for communication, created via **mach_port_allocate()**
 - Send and receive are flexible, for example four options if mailbox full:
 - Wait indefinitely
 - Wait at most n milliseconds
 - Return immediately
 - Temporarily cache a message

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

62



Mach Messages

```
#include<mach/mach.h>

struct message {
    mach_msg_header_t header;
    int data;
};

mach_port_t client;
mach_port_t server;
```

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

63



Mach Message Passing - Client

/ Client Code */*

```
struct message message;

// construct the header
message.header.msgh_size = sizeof(message);
message.header.msgh_remote_port = server;
message.header.msgh_local_port = client;

// send the message
mach_msg(&message.header, // message header
        MACH_SEND_MSG, // sending a message
        sizeof(message), // size of message sent
        0, // maximum size of received message - unnecessary
        MACH_PORT_NULL, // name of receive port - unnecessary
        MACH_MSG_TIMEOUT_NONE, // no time outs
        MACH_PORT_NULL // no notify port
);
```

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

64



Mach Message Passing - Server

`/* Server Code */`

```
struct message message;

// receive the message
mach_msg(&message.header, // message header
MACH_RCV_MSG, // sending a message
0, // size of message sent
sizeof(message), // maximum size of received message
server, // name of receive port
MACH_MSG_TIMEOUT_NONE, // no time outs
MACH_PORT_NULL // no notify port
);
```

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

65



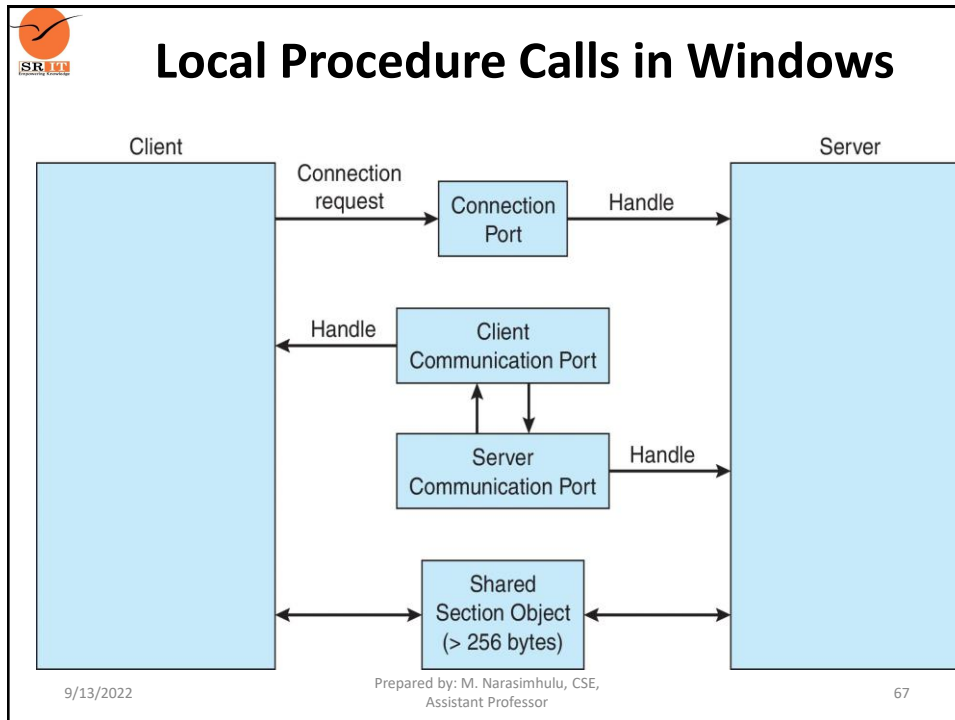
Examples of IPC Systems – Windows

- Message-passing centric via **advanced local procedure call (LPC)** facility
 - Only works between processes on the same system
 - Uses ports (like mailboxes) to establish and maintain communication channels
 - Communication works as follows:
 - The client opens a handle to the subsystem's **connection port** object.
 - The client sends a connection request.
 - The server creates two private **communication ports** and returns the handle to one of them to the client.
 - The client and server use the corresponding port handle to send messages or callbacks and to listen for replies.

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

66



Pipes


- Acts as a conduit allowing two processes to communicate
- Issues:
 - Is communication unidirectional or bidirectional?
 - In the case of two-way communication, is it half or full-duplex?
 - Must there exist a relationship (i.e., **parent-child**) between the communicating processes?
 - Can the pipes be used over a network?
- **Ordinary pipes** – cannot be accessed from outside the process that created it. Typically, a parent process creates a pipe and uses it to communicate with a child process that it created.
- **Named pipes** – can be accessed without a parent-child relationship.

9/13/2022 Prepared by: M. Narasimhulu, CSE, Assistant Professor 68



Ordinary Pipes

- Ordinary Pipes allow communication in standard producer-consumer style
- Producer writes to one end (the **write-end** of the pipe)
- Consumer reads from the other end (the **read-end** of the pipe)
- Ordinary pipes are therefore unidirectional
- Require parent-child relationship between communicating processes



- Windows calls these **anonymous pipes**

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

69



Named Pipes

- Named Pipes are more powerful than ordinary pipes
- Communication is bidirectional
- No parent-child relationship is necessary between the communicating processes
- Several processes can use the named pipe for communication
- Provided on both UNIX and Windows systems

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

70



Communication in Client-Server systems

Narasimhulu M_{M. Tech.}

Assistant Professor

Department of Computer Science & Engineering



Communications in Client-Server Systems

- Sockets
- Remote Procedure Calls



Sockets

- A **socket** is defined as an endpoint for communication
- Concatenation of IP address and **port**
 - port is a number included at start of message packet to differentiate network services on a host
- The socket **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**
- Communication consists between a pair of sockets
- All ports below 1024 are **well known**, used for standard services
- Special IP address 127.0.0.1 (**loopback**) to refer to system on which process is running


9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

73



Socket Communication



9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

74



Sockets in Java

- Three types of sockets
 - **Connection-oriented (TCP)**
 - **Connectionless (UDP)**
 - **MulticastSocket** class– data can be sent to multiple recipients
- Consider this “Date” server in Java:

```
import java.net.*;
import java.io.*;

public class DateServer
{
    public static void main(String[] args) {
        try {
            ServerSocket sock = new ServerSocket(6013);

            /* now listen for connections */
            while (true) {
                Socket client = sock.accept();

                PrintWriter pout = new
                    PrintWriter(client.getOutputStream(), true);

                /* write the Date to the socket */
                pout.println(new java.util.Date().toString());

                /* close the socket and resume */
                /* listening for connections */
                client.close();
            }
        } catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

75



Sockets in Java

The equivalent Date client

```
import java.net.*;
import java.io.*;

public class DateClient
{
    public static void main(String[] args) {
        try {
            /* make connection to server socket */
            Socket sock = new Socket("127.0.0.1", 6013);

            InputStream in = sock.getInputStream();
            BufferedReader bin = new
                BufferedReader(new InputStreamReader(in));

            /* read the date from the socket */
            String line;
            while ( (line = bin.readLine()) != null)
                System.out.println(line);

            /* close the socket connection*/
            sock.close();
        } catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

76



Remote Procedure Calls

- Remote procedure call (RPC) abstracts procedure calls between processes on networked systems
 - Again, uses ports for service differentiation
- **Stubs** – client-side proxy for the actual procedure on the server
- The client-side stub locates the server and **marshalls** the parameters
- The server-side stub receives this message, unpacks the marshalled parameters, and performs the procedure on the server
- On Windows, stub code compile from specification written in **Microsoft Interface Definition Language (MIDL)**

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

77



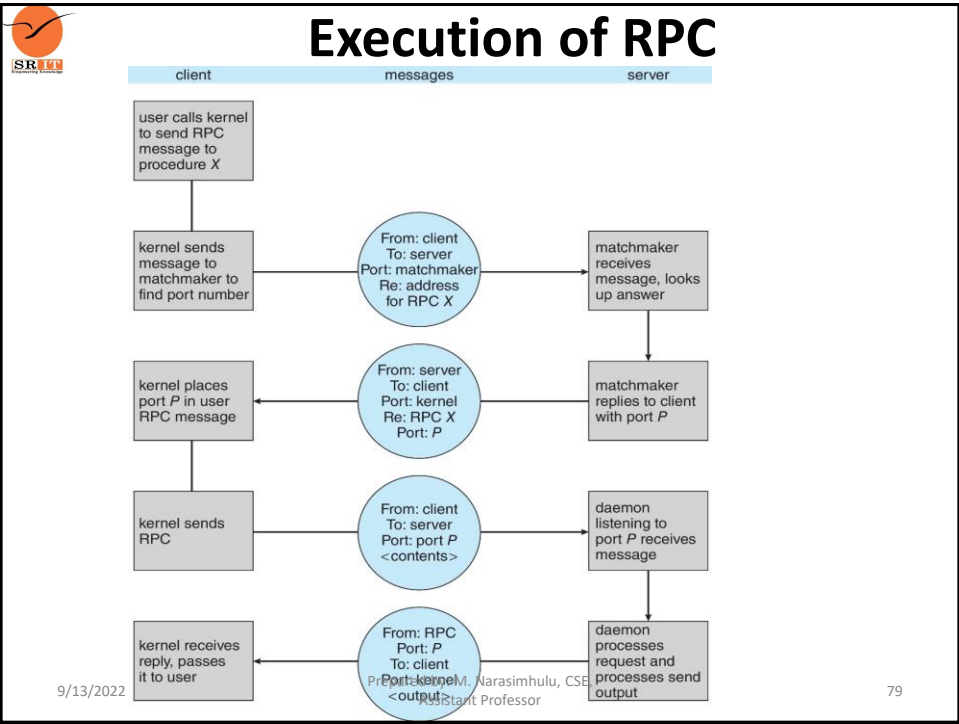
Remote Procedure Calls (Cont.)


- Data representation handled via **External Data Representation (XDL)** format to account for different architectures
 - **Big-endian** and **little-endian**
- Remote communication has more failure scenarios than local
 - Messages can be delivered *exactly once* rather than *at most once*
- OS typically provides a rendezvous (or **matchmaker**) service to connect client and server

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

78





CPU scheduling

Narasimhulu M_{M. Tech.}
Assistant Professor
Department of Computer Science & Engineering



Outline

- Basic Concepts
- Scheduling Criteria
- Scheduling Algorithms

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

81




Objectives

- Describe various CPU scheduling algorithms
- Assess CPU scheduling algorithms based on scheduling criteria
- Explain the issues related to multiprocessor and multicore scheduling
- Describe various real-time scheduling algorithms
- Describe the scheduling algorithms used in the Windows, Linux, and Solaris operating systems
- Apply modeling and simulations to evaluate CPU scheduling algorithms

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

82



Basic Concepts

- Maximum CPU utilization obtained with multi-programming
- CPU-I/O Burst Cycle – Process execution consists of a **cycle** of CPU execution and I/O wait
- CPU burst** followed by **I/O burst**
- CPU burst distribution is of main concern

⋮

load store
add store
read from file

wait for I/O

store increment
index
write to file

wait for I/O

load store
add store
read from file

wait for I/O

⋮

} CPU burst

} I/O burst

} CPU burst

} I/O burst


} CPU burst

} I/O burst

9/13/2022

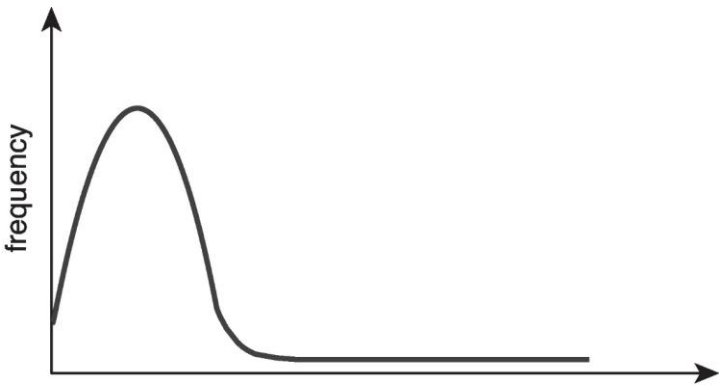
Prepared by: M. Narasimhulu, CSE,
Assistant Professor

83



Histogram of CPU-burst Times

- Large number of short bursts
- Small number of longer bursts
- Histogram



9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

84



CPU Scheduler

- The **CPU scheduler** selects from among the processes in ready queue, and allocates a CPU core to one of them
 - The ready queue may be ordered in various ways
- CPU scheduling decisions may take place when a process:
 1. Switches from running to waiting state
 2. Switches from running to ready state
 3. Switches from waiting to ready
 4. Terminates
- For situations 1 and 4, there is no choice in terms of scheduling. A new process (if one exists in the ready queue) must be selected for execution.
- For situations 2 and 3, however, there is a choice.

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

85



Preemptive and Nonpreemptive Scheduling

- When scheduling takes place only under circumstances 1 and 4, the scheduling scheme is **nonpreemptive**.
- Otherwise, it is **preemptive**.
- Under Nonpreemptive scheduling, once the CPU has been allocated to a process, the process keeps the CPU until it releases it either by terminating or by switching to the waiting state.
 - What is the potential problem?
- Virtually all modern operating systems including Windows, MacOS, Linux, and UNIX use preemptive scheduling algorithms.

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

86



Preemptive Scheduling and Race Conditions

- Preemptive scheduling can result in race conditions when data are shared among several processes.
- Consider the case of two processes that share data. While one process is updating the data, it is preempted so that the second process can run. The second process then tries to read the data, which are in an inconsistent state.
 - We saw this in the bounded buffer example
- This issue will be explored in detail in Chapter 6.

9/13/2022


Prepared by: M. Narasimhulu, CSE,
Assistant Professor

87



Dispatcher

- Dispatcher module gives control of the CPU to the process selected by the CPU scheduler; this involves:
 - Switching context
 - Switching to user mode
 - Jumping to the proper location in the user program to restart that program
- **Dispatch latency** – time it takes for the dispatcher to stop one process and start another running



9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

88



Scheduling Criteria

Narasimhulu M. M. Tech.

Assistant Professor

Department of Computer Science & Engineering



Scheduling Criteria

- **CPU utilization** – keep the CPU as busy as possible
- **Throughput** – # of processes that complete their execution per time unit
- **Turnaround time** – amount of time to execute a particular process
- **Waiting time** – amount of time a process has been waiting in the ready queue
- **Response time** – amount of time it takes from when a request was submitted until the first response is produced

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

90



Optimization Criteria for Scheduling Algorithms

- Max CPU utilization
- Max throughput
- Min turnaround time
- Min waiting time
- Min response time

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

91



Scheduling Algorithms

Narasimhulu M_{M. Tech.}

Assistant Professor

Department of Computer Science & Engineering



First- Come, First-Served (FCFS) Scheduling

- Example with 3 processes

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

- Suppose that the processes arrive in the order: P_1, P_2, P_3
The Gantt Chart for the above schedule is:



- Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- Average waiting time: $(0 + 24 + 27)/3 = 17$

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

93



FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order:

$$P_2, P_3, P_1$$

- The Gantt chart for the schedule is:




- Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
- Average waiting time: $(6 + 0 + 3)/3 = 3$
- Much better than previous case
- Convoy effect** - short process behind long process
 - Consider one CPU-bound and many I/O-bound processes

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

94




Shortest-Job-First (SJF) Scheduling

- Associate with each process the length of its next CPU burst
 - Use these lengths to schedule the process with the shortest time
- SJF is optimal – gives minimum average waiting time for a given set of processes
- How do we determine the length of the next CPU burst?
 - Could ask the user
 - Estimate

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

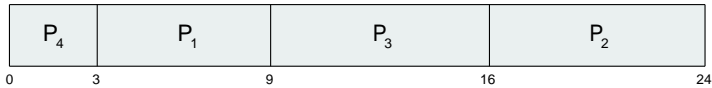
95



Example of SJF

<u>Process</u>	<u>Burst Time</u>
P_1	6
P_2	8
P_3	7
P_4	3

- SJF scheduling chart




0 3 9 16 24

- Average waiting time = $(3 + 16 + 9 + 0) / 4 = 7$

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

96




Determining Length of Next CPU Burst

- Can only estimate the length – should be similar to the previous one
 - Then pick process with shortest predicted next CPU burst
- Can be done by using the length of previous CPU bursts, using exponential averaging
 - t_n = actual length of n^{th} CPU burst
 - τ_{n+1} = predicted value for the next CPU burst
 - $\alpha, 0 \leq \alpha \leq 1$
 - Define :
$$\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n.$$
- Commonly, α set to $\frac{1}{2}$

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

97





Examples of Exponential Averaging

- $\alpha = 0$
 - $\tau_{n+1} = \tau_n$
 - Recent history does not count
- $\alpha = 1$
 - $\tau_{n+1} = \alpha t_n$
 - Only the actual last CPU burst counts
- If we expand the formula, we get:

$$\tau_{n+1} = \alpha t_n + (1 - \alpha)\alpha t_{n-1} + \dots$$

$$+ (1 - \alpha)^j \alpha t_{n-j} + \dots$$

$$+ (1 - \alpha)^{n+1} \tau_0$$
- Since both α and $(1 - \alpha)$ are less than or equal to 1, each successive term has less weight than its predecessor

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

99




Shortest Remaining Time First Scheduling

- Preemptive version of SJN
- Whenever a new process arrives in the ready queue, the decision on which process to schedule next is redone using the SJN algorithm.
- Is SRT more “optimal” than SJN in terms of the minimum average waiting time for a given set of processes?

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

100

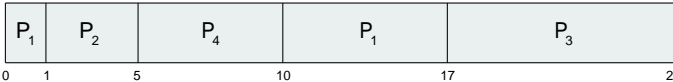


Example of Shortest-remaining-time-first

- Now we add the concepts of varying arrival times and preemption to the analysis

Process	Arrival Time	Burst Time
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

- Preemptive SJF Gantt Chart




- Average waiting time = $[(10-1)+(1-1)+(17-2)+(5-3)]/4 = 26/4 = 6.5$

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

101




Round Robin (RR)

- Each process gets a small unit of CPU time (**time quantum q**), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in chunks of at most q time units at once. No process waits more than $(n-1)q$ time units.
- Timer interrupts every quantum to schedule next process
- Performance
 - q large \Rightarrow FIFO
 - q small $\Rightarrow q$ must be large with respect to context switch, otherwise overhead is too high

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

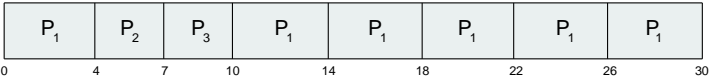
102



Example of RR with Time Quantum = 4

Process	Burst Time
P_1	24
P_2	3
P_3	3


- The Gantt chart is:






- Typically, higher average turnaround than SJF, but better **response**
- q should be large compared to context switch time
 - q usually 10 milliseconds to 100 milliseconds,
 - Context switch < 10 microseconds

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

103




Time Quantum and Context Switch Time

process time = 10	quantum	context switches
 0 10	12	0
 0 6 10	6	1
 0 1 2 3 4 5 6 7 8 9 10	1	9

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor


104



Priority Scheduling

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (usually, smallest integer \equiv highest priority)
- Two schemes:
 - Preemptive
 - Nonpreemptive
- Problem \equiv **Starvation** – low priority processes may never execute
- Solution \equiv **Aging** – as time progresses increase the priority of the process
- Note: SJF is priority scheduling where priority is the inverse of predicted next CPU burst time

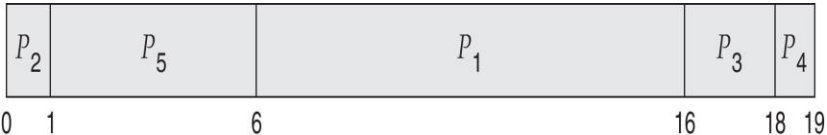
Prepared by: M. Narasimhulu, CSE,
Assistant Professor



Example of Priority Scheduling

Process	Burst Time	Priority
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2

- Priority scheduling Gantt Chart




- Average waiting time = 8.2

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

107

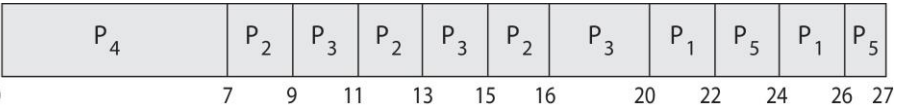


Priority Scheduling w/ Round-Robin

- Run the process with the highest priority. Processes with the same priority run round-robin
- Example:

Process	Burst Time	Priority
P_1	4	3
P_2	5	2
P_3	8	2
P_4	7	1
P_5	3	3


- Gantt Chart with time quantum = 2



9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

108



Multilevel Queue

- The ready queue consists of multiple queues
- Example:
 - Priority scheduling, where each priority has its separate queue.
 - Schedule the process in the highest-priority queue!

priority = 0

T₀ T₁ T₂ T₃ T₄

priority = 1

T₅ T₆ T₇

priority = 2

T₈ T₉ T₁₀ T₁₁

⋮


priority = n

T_x T_y T_z

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

109



Multilevel Queue

- Prioritization based upon process type

highest priority

→

real-time processes

→

→

system processes

→

→

interactive processes

→

→

batch processes

→

lowest priority

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

110



Multilevel Feedback Queue

- A process can move between the various queues.
- Multilevel-feedback-queue scheduler defined by the following parameters:
 - Number of queues
 - Scheduling algorithms for each queue
 - Method used to determine when to upgrade a process
 - Method used to determine when to demote a process
 - Method used to determine which queue a process will enter when that process needs service
- Aging can be implemented using multilevel feedback queue

9/13/2022


Prepared by: M. Narasimhulu, CSE,
Assistant Professor

111



Example of Multilevel Feedback Queue

- Three queues:
 - Q_0 – RR with time quantum 8 milliseconds
 - Q_1 – RR time quantum 16 milliseconds
 - Q_2 – FCFS
- Scheduling
 - A new process enters queue Q_0 which is served in RR
 - When it gains CPU, the process receives 8 milliseconds
 - If it does not finish in 8 milliseconds, the process is moved to queue Q_1
 - At Q_1 job is again served in RR and receives 16 additional milliseconds
 - If it still does not complete, it is preempted and moved to queue Q_2



9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

112



Threads

Narasimhulu M_{M. Tech.}

Assistant Professor

Department of Computer Science & Engineering



Thread Scheduling

- Distinction between user-level and kernel-level threads
- When threads supported, threads scheduled, not processes
- Many-to-one and many-to-many models, thread library schedules user-level threads to run on LWP
 - Known as **process-contention scope (PCS)** since scheduling competition is within the process
 - Typically done via priority set by programmer
- Kernel thread scheduled onto available CPU is **system-contention scope (SCS)** – competition among all threads in system



Pthread Scheduling

- API allows specifying either PCS or SCS during thread creation
 - PTHREAD_SCOPE_PROCESS schedules threads using PCS scheduling
 - PTHREAD_SCOPE_SYSTEM schedules threads using SCS scheduling
- Can be limited by OS – Linux and macOS only allow PTHREAD_SCOPE_SYSTEM

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

115



Pthread Scheduling API

```
#include <pthread.h>
#include <stdio.h>
#define NUM_THREADS 5
int main(int argc, char *argv[]) {
    int i, scope;
    pthread_t tid[NUM_THREADS];
    pthread_attr_t attr;
    /* get the default attributes */
    pthread_attr_init(&attr);
    /* first inquire on the current scope */
    if (pthread_attr_getscope(&attr, &scope) != 0)
        fprintf(stderr, "Unable to get scheduling scope\n");
    else {
        if (scope == PTHREAD_SCOPE_PROCESS)
            printf("PTHREAD_SCOPE_PROCESS");
        else if (scope == PTHREAD_SCOPE_SYSTEM)
            printf("PTHREAD_SCOPE_SYSTEM");
        else
            fprintf(stderr, "Illegal scope value.\n");
    }
}
```

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

116



Pthread Scheduling API

```

/* set the scheduling algorithm to PCS or SCS */
pthread_attr_setscope(&attr, PTHREAD_SCOPE_SYSTEM);
/* create the threads */
for (i = 0; i < NUM_THREADS; i++)
    pthread_create(&tid[i], &attr, runner, NULL);
/* now join on each thread */
for (i = 0; i < NUM_THREADS; i++)
    pthread_join(tid[i], NULL);
}
/* Each thread will begin control in this function */
void *runner(void *param)
{
    /* do some work ... */
    pthread_exit(0);
}

```

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

117



END of Chapter - 1

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

118



Chapter 2

Process Synchronization

Narasimhulu M_{M. Tech.}
Assistant Professor
Department of Computer Science & Engineering



Critical-Section Problem

Narasimhulu M_{M. Tech.}
Assistant Professor
Department of Computer Science & Engineering



Background

- Processes can execute concurrently
 - May be interrupted at any time, partially completing execution
- Concurrent access to shared data may result in data inconsistency
- Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes
- We illustrated in chapter 4 the problem when we considered the Bounded Buffer problem with use of a counter that is updated concurrently by the producer and consumer,. Which lead to race condition.

9/13/2022

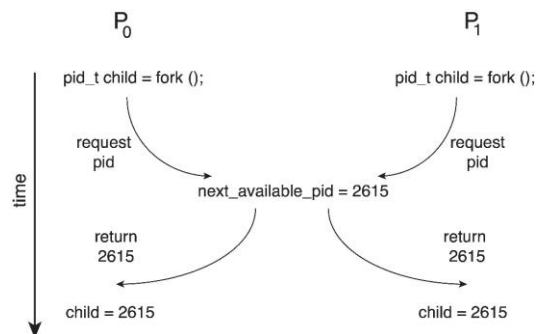
Prepared by: M. Narasimhulu, CSE,
Assistant Professor

121



Race Condition

- Processes P_0 and P_1 are creating child processes using the `fork()` system call
- Race condition on kernel variable `next_available_pid` which represents the next available process identifier (pid)



- Unless there is a mechanism to prevent P_0 and P_1 from accessing the variable `next_available_pid` the same pid could be assigned to two different processes!

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

122



Critical Section Problem

- Consider system of n processes $\{p_0, p_1, \dots, p_{n-1}\}$
- Each process has **critical section** segment of code
 - Process may be changing common variables, updating table, writing file, etc.
 - When one process in critical section, no other may be in its critical section
- **Critical section problem** is to design protocol to solve this
- Each process must ask permission to enter critical section in **entry section**, may follow critical section with **exit section**, then **remainder section**

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

123



Critical Section

- General structure of process P_i

```

while (true) {
    entry section
    critical section
    exit section
    remainder section
}
  
```

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

124



Critical-Section Problem (Cont.)

Requirements for solution to critical-section problem

1. **Mutual Exclusion** - If process P_i is executing in its critical section, then no other processes can be executing in their critical sections
2. **Progress** - If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the selection of the process that will enter the critical section next cannot be postponed indefinitely
3. **Bounded Waiting** - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted
 - Assume that each process executes at a nonzero speed
 - No assumption concerning **relative speed** of the n processes

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

125



Interrupt-based Solution

- Entry section: disable interrupts
- Exit section: enable interrupts
- Will this solve the problem?
 - What if the critical section is code that runs for an hour?
 - Can some processes starve – never enter their critical section.
 - What if there are two CPUs?

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

126



Software Solution 1

- Two process solution
- Assume that the `load` and `store` machine-language instructions are atomic; that is, cannot be interrupted
- The two processes share one variable:
 - `int turn;`
- The variable `turn` indicates whose turn it is to enter the critical section
- initially, the value of `turn` is set to i

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

127



Algorithm for Process P_i

```
while (true){
```

```
    while (turn == j);
```

```
    /* critical section */
```

```
    turn = j;
```

```
    /* remainder section */
```

```
}
```

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

128



Correctness of the Software Solution

- Mutual exclusion is preserved

P_i enters critical section only if:

$$\text{turn} = i$$

and turn cannot be both 0 and 1 at the same time

- What about the Progress requirement?
- What about the Bounded-waiting requirement?

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

129



Petersons Solution

Narasimhulu M. M. Tech.

Assistant Professor

Department of Computer Science & Engineering



Peterson's Solution

- Two process solution
- Assume that the `load` and `store` machine-language instructions are atomic; that is, cannot be interrupted
- The two processes share two variables:
 - `int turn;`
 - `boolean flag[2]`
- The variable `turn` indicates whose turn it is to enter the critical section
- The `flag` array is used to indicate if a process is ready to enter the critical section.

9/13/2022

`flag[i] = true` implies that process P_i is ready!

Prepared by: M. Narasimhulu, CSE,
Assistant Professor



Algorithm for Process P_i

```
while (true){
    flag[i] = true;
    turn = j;
    while (flag[j] && turn == j)
        ;

    /* critical section */

    flag[i] = false;

    /* remainder section */
}
```

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

132



Correctness of Peterson's Solution

- Provable that the three CS requirement are met:
 1. Mutual exclusion is preserved
 - P_i enters CS only if:
 - either `flag[j] = false` Or `turn = i`
 2. Progress requirement is satisfied
 3. Bounded-waiting requirement is met

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

133



Peterson's Solution and Modern Architecture

- Although useful for demonstrating an algorithm, Peterson's Solution is not guaranteed to work on modern architectures.
 - To improve performance, processors and/or compilers may reorder operations that have no dependencies
- Understanding why it will not work is useful for better understanding race conditions.
- For single-threaded this is ok as the result will always be the same.
- For multithreaded the reordering may produce inconsistent or unexpected results!

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

134



Modern Architecture Example

- Two threads share the data:

```
boolean flag = false;  
int x = 0;
```
- Thread 1 performs

```
while (!flag);  
print x
```
- Thread 2 performs

```
x = 100;  
flag = true
```
- What is the expected output?

100

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

135



Modern Architecture Example (Cont.)

- However, since the variables `flag` and `x` are independent of each other, the instructions:

```
flag = true;  
x = 100;
```

for Thread 2 may be reordered

- If this occurs, the output may be 0!

9/13/2022


Prepared by: M. Narasimhulu, CSE,
Assistant Professor

136



Peterson's Solution Revisited

- The effects of instruction reordering in Peterson's Solution



- This allows both processes to be in their critical section at the same time!
- To ensure that Peterson's solution will work correctly on modern computer architecture we must use **Memory Barrier**.

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

137



Synchronization Hardware

Narasimhulu M. *M. Tech.*

Assistant Professor

Department of Computer Science & Engineering



Memory Barrier

- **Memory model** are the memory guarantees a computer architecture makes to application programs.
- Memory models may be either:
 - **Strongly ordered** – where a memory modification of one processor is immediately visible to all other processors.
 - **Weakly ordered** – where a memory modification of one processor may not be immediately visible to all other processors.
- A **memory barrier** is an instruction that forces any change in memory to be propagated (made visible) to all other processors.

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

139



Memory Barrier Instructions

- When a memory barrier instruction is performed, the system ensures that all loads and stores are completed before any subsequent load or store operations are performed.
- Therefore, even if instructions were reordered, the memory barrier ensures that the store operations are completed in memory and visible to other processors before future load or store operations are performed.

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

140



Memory Barrier Example

- Returning to the example of slides 6.17 - 6.18
- We could add a memory barrier to the following instructions to ensure Thread 1 outputs 100:
- Thread 1 now performs


```
while (!flag)
    memory_barrier();
print x
```
- Thread 2 now performs


```
x = 100;
memory_barrier();
flag = true
```
- For Thread 1 we are guaranteed that the value of `flag` is loaded before the value of `x`.
- For Thread 2 we ensure that the assignment to `x` occurs before the assignment `flag`.

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

141



Synchronization Hardware

- Many systems provide hardware support for implementing the critical section code.
- Uniprocessors – could disable interrupts
 - Currently running code would execute without preemption
 - Generally too inefficient on multiprocessor systems
 - Operating systems using this not broadly scalable
- We will look at three forms of hardware support:
 1. Hardware instructions
 2. Atomic variables

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

142



Hardware Instructions

- Special hardware instructions that allow us to either *test-and-modify* the content of a word, or to *swap* the contents of two words atomically (uninterruptedly.)
 - **Test-and-Set** instruction
 - **Compare-and-Swap** instruction

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

143



The test_and_set Instruction

- Definition


```

boolean test_and_set (boolean
*target)
{
    boolean rv = *target;
    *target = true;
    return rv;
}

```
- Properties
 - Executed atomically
 - Returns the original value of passed parameter
 - Set the new value of passed parameter to **true**

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

144



Solution Using test_and_set()

- Shared boolean variable **lock**, initialized to **false**
- Solution:

```
do {
    while (test_and_set(&lock))
        ; /* do nothing */

        /* critical section */

    lock = false;
        /* remainder section */
} while (true);
```

- Does it solve the critical-section problem?

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

145



The compare_and_swap Instruction

- Definition

```
int compare_and_swap(int *value, int expected, int
    new_value){
    int temp = *value;
    if (*value == expected)
        *value = new_value;
    return temp; }
```

- Properties

- Executed atomically
- Returns the original value of passed parameter **value**
- Set the variable **value** the value of the passed parameter **new_value** but only if ***value == expected** is true. That is, the swap takes place only under this condition.

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

146



Solution using compare_and_swap

- Shared integer `lock` initialized to 0;
- Solution:

```
while (true){
    while (compare_and_swap(&lock, 0, 1) != 0)
        ; /* do nothing */

    /* critical section */

    lock = 0;

    /* remainder section */
}
```

- Does it solve the critical-section problem?

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

147



Bounded-waiting with compare-and-swap

```
while (true) {
    waiting[i] = true;
    key = 1;
    while (waiting[i] && key == 1)
        key = compare_and_swap(&lock, 0, 1);
    waiting[i] = false;
    /* critical section */
    j = (i + 1) % n;
    while ((j != i) && !waiting[j])
        j = (j + 1) % n;
    if (j == i)
        lock = 0;
    else
        waiting[j] = false;
    /* remainder section */
}
```

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

148



Atomic Variables

- Typically, instructions such as compare-and-swap are used as building blocks for other synchronization tools.
- One tool is an **atomic variable** that provides *atomic* (uninterruptible) updates on basic data types such as integers and booleans.
- For example:
 - Let `sequence` be an atomic variable
 - Let `increment()` be operation on the atomic variable `sequence`
 - The Command:

```
increment(&sequence);
```

ensures `sequence` is incremented without interruption:

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

149



Atomic Variables

- The `increment()` function can be implemented as follows:

```
void increment(atomic_int *v)
{
    int temp;
    do {
        temp = *v;
    }
    while (temp !=
(compare_and_swap(v, temp, temp+1)) );
}
```

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

150



Mutex Locks

Narasimhulu M M. Tech.

Assistant Professor

Department of Computer Science & Engineering




Mutex Locks

- Previous solutions are complicated and generally inaccessible to application programmers
- OS designers build software tools to solve critical section problem
- Simplest is `mutex` lock
 - Boolean variable indicating if lock is available or not
- Protect a critical section by
 - First `acquire()` a lock
 - Then `release()` the lock
- Calls to `acquire()` and `release()` must be **atomic**
 - Usually implemented via hardware atomic instructions such as compare-and-swap.
- But this solution requires **busy waiting**
 - This lock therefore called a **spinlock**.

9/13/2022

Mr. M. Narasimhulu, CSE,
Assistant Professor

152




Solution to CS Problem Using Mutex Locks

```
while (true) {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
}
```

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

153



Semaphores

Narasimhulu M_{M. Tech.}
Assistant Professor
Department of Computer Science & Engineering



Semaphore

- Synchronization tool that provides more sophisticated ways (than Mutex locks) for processes to synchronize their activities.
- Semaphore S – integer variable
- Can only be accessed via two indivisible (atomic) operations
 - `wait()` and `signal()`
 - Originally called (proberen) $P()$ and (verhogen) $V()$
- Definition of the `wait()` operation

```
wait(S) {
    while (S <= 0)
        ; // busy wait
    S--;
}
```

- Definition of the `signal()` operation

```
signal(S) {
    S++;
}
```

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

155



Semaphore (Cont.)

- **Counting semaphore** – integer value can range over an unrestricted domain
- **Binary semaphore** – integer value can range only between 0 and 1
 - Same as a **mutex lock**
- Can implement a counting semaphore S as a binary semaphore
- With semaphores we can solve various synchronization problems

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

156



Semaphore Usage Example

- Solution to the CS Problem
 - Create a semaphore “**mutex**” initialized to 1


```
wait(mutex);
CS
signal(mutex);
```
- Consider P_1 and P_2 that with two statements S_1 and S_2 and the requirement that S_1 to happen before S_2
 - Create a semaphore “**synch**” initialized to 0


```
P1:
    S1;
    signal(synch);
P2:
    wait(synch);
    S2;
```

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

157



Semaphore Implementation

- Must guarantee that no two processes can execute the **wait()** and **signal()** on the same semaphore at the same time
- Thus, the implementation becomes the critical section problem where the **wait** and **signal** code are placed in the critical section
- Could now have **busy waiting** in critical section implementation
 - But implementation code is short
 - Little busy waiting if critical section rarely occupied
- Note that applications may spend lots of time in critical sections and therefore this is not a good solution

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

158



Semaphore Implementation with no Busy waiting

- With each semaphore there is an associated waiting queue
- Each entry in a waiting queue has two data items:
 - Value (of type integer)
 - Pointer to next record in the list
- Two operations:
 - **block** – place the process invoking the operation on the appropriate waiting queue
 - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

159



Implementation with no Busy waiting (Cont.)

- Waiting queue

```
typedef struct {
    int value;
    struct process *list;
} semaphore;
```

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

160



Implementation with no Busy waiting (Cont.)

```
wait(semaphore *S) {
    S->value--;
    if (S->value < 0) {
        add this process to S->list;
        block();
    }
}

signal(semaphore *S) {
    S->value++;
    if (S->value <= 0) {
        remove a process P from S->list;
        wakeup(P);
    }
}
```

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

161



Problems with Semaphores

- Incorrect use of semaphore operations:
 - `signal(mutex) ... wait(mutex)`
 - `wait(mutex) ... wait(mutex)`
 - Omitting of `wait (mutex)` and/or `signal (mutex)`
- These – and others – are examples of what can occur when semaphores and other synchronization tools are used incorrectly.

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

162



Monitors

Narasimhulu M. *M. Tech.*

Assistant Professor

Department of Computer Science & Engineering



Monitors

- A high-level abstraction that provides a convenient and effective mechanism for process synchronization
- *Abstract data type*, internal variables only accessible by code within the procedure
- Only one process may be active within the monitor at a time
- Pseudocode syntax of a monitor:

```
monitor monitor-name
{
    // shared variable declarations
    procedure P1 (...) { ... }

    procedure      P2      (...)      {      ...      }

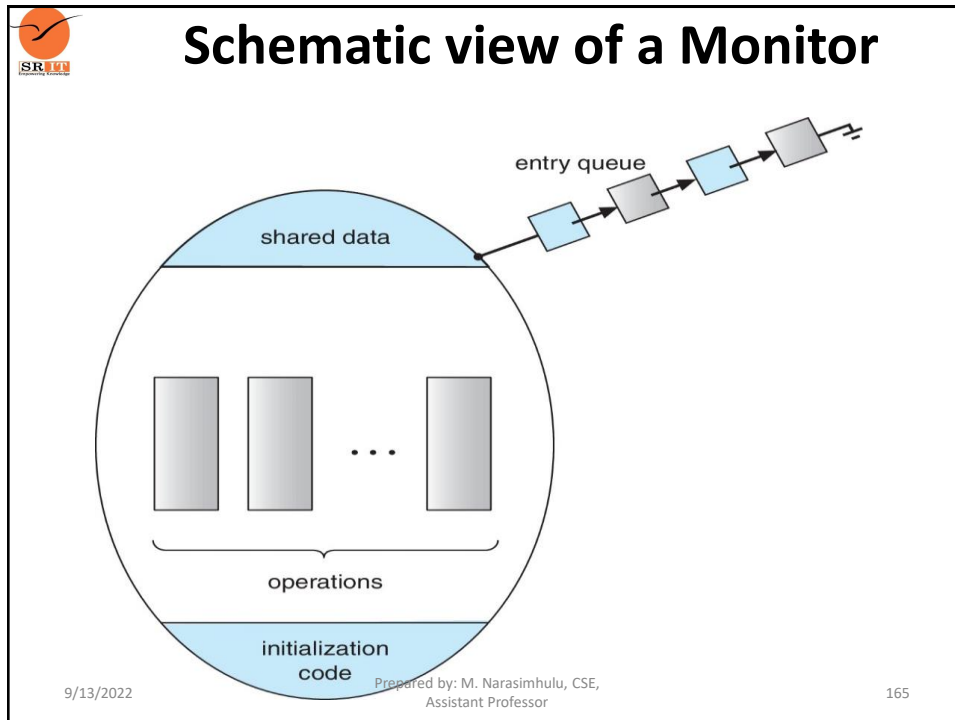
    procedure Pn (...) {.....}

    initialization code (...) { ... }
}
```

9/13/2022 }

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

164



Monitor Implementation Using Semaphores

- Variables


```
semaphore mutex
mutex = 1
```
- Each procedure **P** is replaced by


```
wait(mutex) ;
...
body of P;
...
signal(mutex) ;
```
- Mutual exclusion within a monitor is ensured

9/13/2022 Prepared by: M. Narasimhulu, CSE,
Assistant Professor 166



Condition Variables

- **condition x, y ;**
- Two operations are allowed on a condition variable:
 - **$x.wait()$** – a process that invokes the operation is suspended until **$x.signal()$**
 - **$x.signal()$** – resumes one of processes (if any) that invoked **$x.wait()$**
 - If no **$x.wait()$** on the variable, then it has no effect on the variable


9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

167



Monitor with Condition Variables



9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

168



Usage of Condition Variable Example

- Consider P_1 and P_2 that need to execute two statements S_1 and S_2 and the requirement that S_1 to happen before S_2
 - Create a monitor with two procedures F_1 and F_2 that are invoked by P_1 and P_2 respectively
 - One condition variable "x" initialized to 0
 - One Boolean variable "done"
 - F1:**

```
S1;
done = true;
x.signal();
```
 - F2:**

```
if done = false
    x.wait();
S2;
```

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

169



Monitor Implementation Using Semaphores

Variables

```
semaphore mutex; // (initially = 1)
semaphore next;  // (initially = 0)
int next_count = 0; // number of processes waiting
                    inside the monitor
```

Each function P will be replaced by

```
wait(mutex);
...
body of P;
...
if (next_count > 0)
    signal(next)
else
    signal(mutex);
```

Mutual exclusion within a monitor is ensured

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

170



Implementation – Condition Variables

- For each condition variable x , we have:

```
semaphore x_sem; // (initially = 0)
int x_count = 0;
```

- The operation $x.\text{wait}()$ can be implemented as:

```
x_count++;
if (next_count > 0)
    signal(next);
else
    signal(mutex);
wait(x_sem);
x_count--;
```

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

171



Implementation (Cont.)

- The operation $x.\text{signal}()$ can be implemented as:

```
if (x_count > 0) {
    next_count++;
    signal(x_sem);
    wait(next);
    next_count--;
}
```

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

172



Resuming Processes within a Monitor

- If several processes queued on condition variable `x`, and `x.signal()` is executed, which process should be resumed?
- FCFS frequently not adequate
- Use the **conditional-wait** construct of the form
`x.wait(c)`

where:

- `c` is an integer (called the priority number)
- The process with lowest number (highest priority) is scheduled next

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

173



Single Resource allocation

- Allocate a single resource among competing processes using priority numbers that specifies the maximum time a process plans to use the resource

```
R.acquire(t) ;
...
access the resource;
...

R.release;
```

- Where `R` is an instance of type `ResourceAllocator`

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

174



Single Resource allocation

- Allocate a single resource among competing processes using priority numbers that specifies the maximum time a process plans to use the resource
- The process with the shortest time is allocated the resource first
- Let R is an instance of type `ResourceAllocator` (next slide)
- Access to `ResourceAllocator` is done via:

```
R.acquire(t) ;
    . . .
    access the resource;
    . . .
R.release;
```

- Where t is the maximum time a process plans to use the resource

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

175



A Monitor to Allocate Single Resource

```
monitor ResourceAllocator
{
    boolean busy;
    condition x;
    void acquire(int time) {
        if (busy)
            x.wait(time);
        busy = true;
    }
    void release() {
        busy = false;
        x.signal();
    }
    initialization code() {
        busy = false;
    }
}
```

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

176



Single Resource Monitor (Cont.)

- Usage:


```

      acquire
      ...
      release
      
```
- Incorrect use of monitor operations
 - `release()` ... `acquire()`
 - `acquire()` ... `acquire()`
 - Omitting of `acquire()` and/or `release()`

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

177



Liveness

- Processes may have to wait indefinitely while trying to acquire a synchronization tool such as a mutex lock or semaphore.
- Waiting indefinitely violates the progress and bounded-waiting criteria discussed at the beginning of this chapter.
- **Liveness** refers to a set of properties that a system must satisfy to ensure processes make progress.
- Indefinite waiting is an example of a liveness failure.

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

178



Liveness

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes
- Let S and Q be two semaphores initialized to 1

P_0	P_1
<code>wait(S);</code>	<code>wait(Q);</code>
<code>wait(Q);</code>	<code>wait(S);</code>
<code>...</code>	<code>...</code>
<code>signal(S);</code>	<code>signal(Q);</code>
<code>signal(Q);</code>	<code>signal(S);</code>

- Consider if P_0 executes `wait(S)` and P_1 `wait(Q)`. When P_0 executes `wait(Q)`, it must wait until P_1 executes `signal(Q)`
- However, P_1 is waiting until P_0 execute `signal(S)`.
- Since these `signal()` operations will never be executed, P_0 and P_1 are **deadlocked**.

SRIT

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

179



Liveness

- Other forms of deadlock:
- **Starvation** – indefinite blocking
 - A process may never be removed from the semaphore queue in which it is suspended
- **Priority Inversion** – Scheduling problem when lower-priority process holds a lock needed by higher-priority process
 - Solved via **priority-inheritance protocol**

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

180



END of Chapter - 2

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

181



END of Unit-2

9/13/2022

Prepared by: M. Narasimhulu, CSE,
Assistant Professor

182